

PDF with a Real Embedded Image

This document contains a real raster JPEG image embedded directly inside the PDF. It is intended for testing PDF image extraction with Docling, PyMuPDF, pdfplumber, or another document-processing pipeline.




Image caption: This is an actual pixel-based JPEG image embedded in the PDF, not a vector drawing generated by the PDF layout engine.

1. Why this image is useful for testing

The image contains buildings, trees, a road, sky, and text rendered into the pixels. A PDF parser that supports image extraction should be able to identify an image object inside the PDF rather than treating the diagram as ordinary PDF vector graphics.

2. RAG test content

Document processing pipelines often need to handle mixed content. A PDF may contain normal paragraphs, tables, photographs, screenshots, scanned pages, charts, and captions. For multimodal RAG, extracted images can be passed to a vision-language model while the surrounding text can be indexed for retrieval.

3. Content checklist

Element	Expected content
Text	Normal selectable PDF text
Image	Embedded JPEG raster image

Caption	Text describing the image
Table	Structured rows and columns